Your priorities. Your commitments. Followed through.

A full AI executive assistant with five shared engines, approved plans and server-run follow-up.

Product blueprint & UI/UX specification
Version 2.1 / 31 October 2026

Clean Tech edition / Design-stage specification
Includes five engines, fixed workflows, screen designs, shared context, budgets and release gates.

DOCUMENT GUIDE

From intent to execution

This document consolidates the product documentation suite and expands the screen-by-screen design. Confirmed decisions are separated from proposed defaults and unresolved platform choices.

> **Core loop**
> Clarify → Plan → Approve → Schedule → Daily coach sync → Focus → Follow through → Review → Propose adjustments → Approve.

01 Product description

02 Requirements & acceptance

03 End-to-end workflows

04 Interface architecture

05 Screen design & components

06 Features & integrations

07 Templates & background operation

08 Shared context & AI budgets

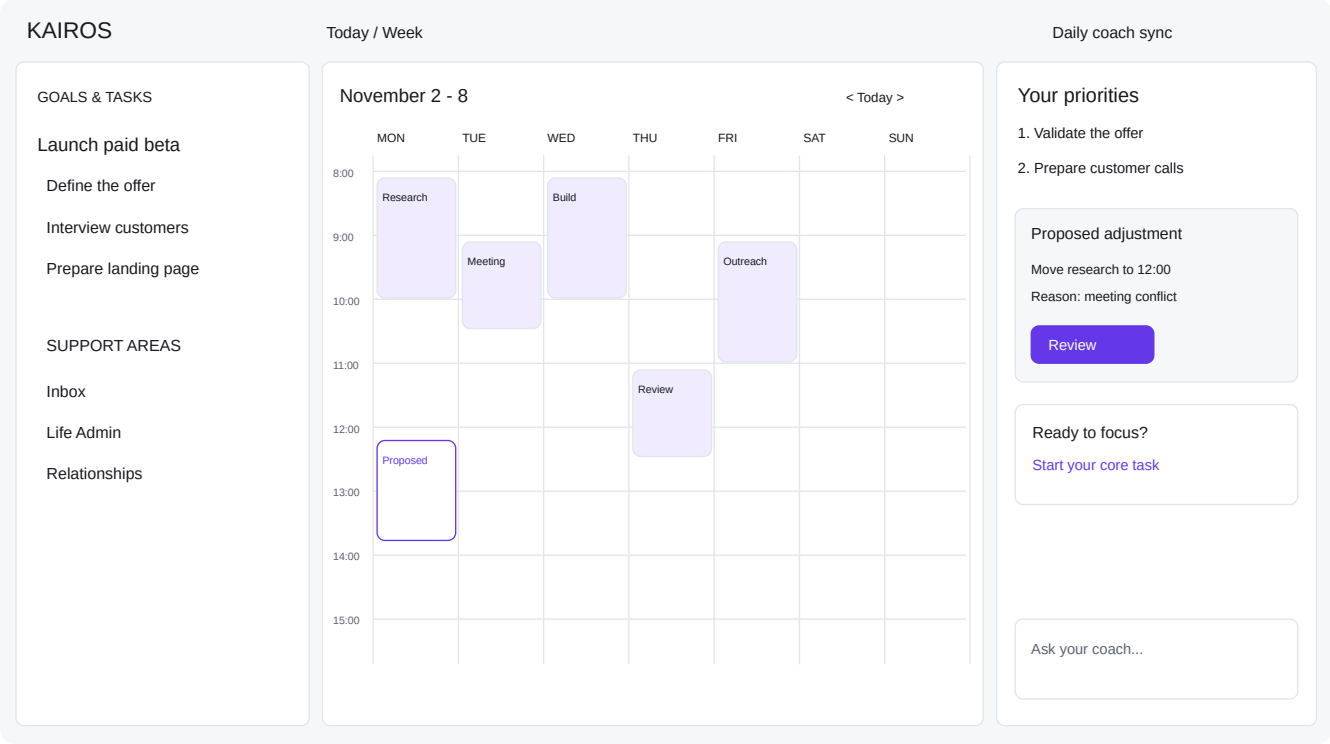
09 Implementation & release gates

10 Decisions & references

Status: design specification. This is not evidence that every feature is implemented and does not authorize deployment. Existing implementation and live-provider acceptance remain separate.

WORKSPACE CONCEPT / S04

A calm calendar at the center



Conceptual desktop layout with the coach open. The wireframe illustrates information hierarchy, not final pixel dimensions.

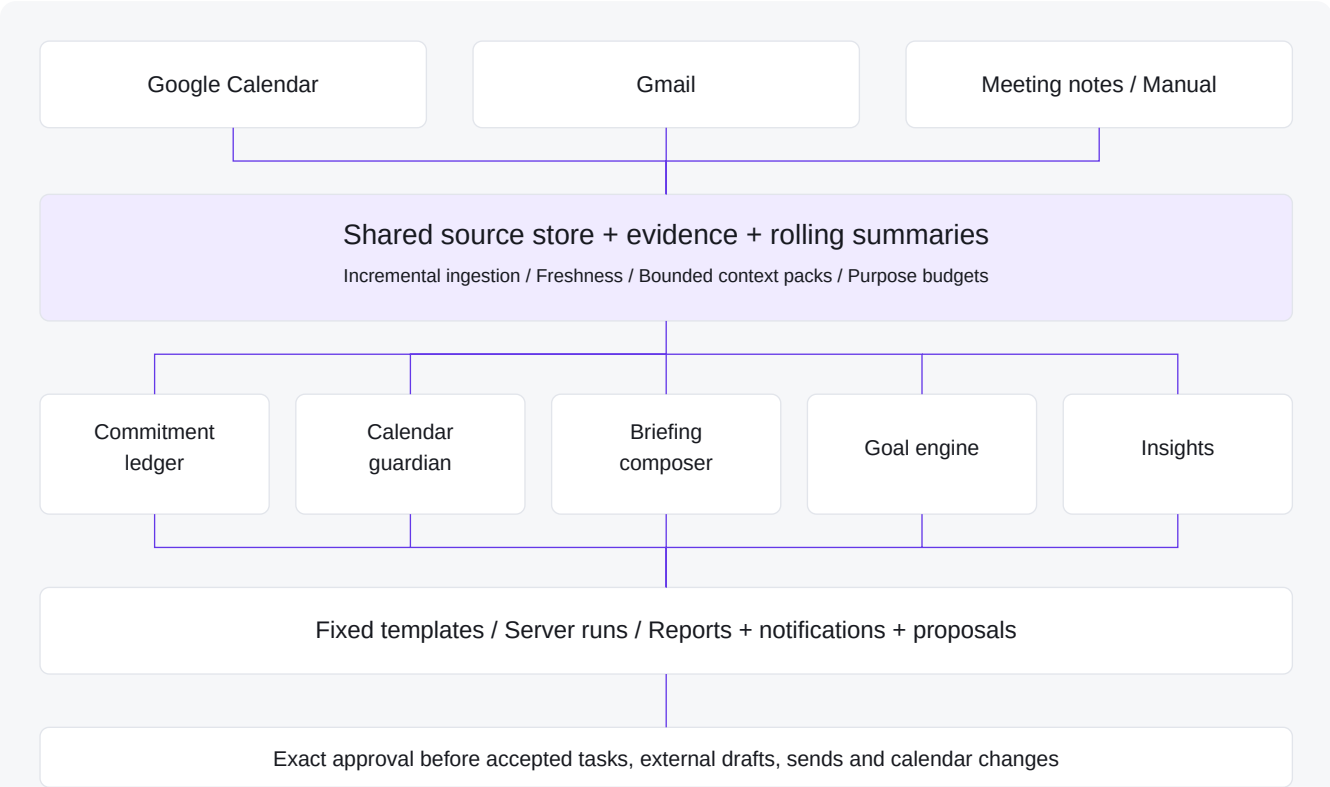
One workspace, three roles

Left: goals, tasks, support areas and quick capture. **Center:** the real week, linked work and proposed adjustments. **Right:** daily coaching, chat, follow-ups and explicit approvals.

> When the coach closes, the calendar expands. On smaller screens, drawers and day/agenda views preserve usable space.

ASSISTANT ARCHITECTURE

Five engines. One shared context.



Logical modules share facts and bounded AI steps. Five engines do not require five continuously running agents or five deployed microservices.

> The first web release includes all five engines. Calendar → meetings → email is internal build order, not a planner-only release.

Product description

Version 2.1 / 31 October 2026. Target product specification; capabilities below are planned unless separately verified in the implementation record.

Purpose

Kairos is a web-first AI executive assistant that turns goals and commitments into an agreed plan, then follows up until each item is completed, cancelled, or deliberately deferred. It serves founders, solopreneurs, freelancers, and team leads in a personal workspace; shared team administration is deferred.

Users already have calendars. Kairos earns its place by keeping a commitment ledger, detecting what needs attention, preparing briefs and drafts, and bringing concrete decisions back for approval. Its calendar-first interface connects this assistant work to available time, goals, and focus.

Product promise

Define what matters. Agree on a realistic plan. Activate follow-up loops. Work with protected focus. Kairos tracks the evidence, prepares the next useful action, and asks before changing the schedule or sending anything.

Five assistant engines

Engine	Responsibility	Authoritative inputs and outputs
Commitment ledger	Track promises, reply debt, waiting items, meeting follow-through and VIP attention	Commitments from permitted email, supplied meeting notes and manual entry; linked tasks, resolution evidence and reviewable follow-ups
Calendar guardian	Detect conflicts and overload, protect/recover focus time, propose freed-time use and travel buffers	Synchronized calendar events and focus reservations; findings and proposed calendar changes
Briefing composer	Prepare daily briefs, week-ahead preparation, meeting dossiers and Friday wraps	Reads shared engine facts and summaries; produces a coherent brief without creating competing facts
Goal engine	Support seasons, goal clarification, GPS-style planning, weekly plans, trajectory and velocity	Goals, tasks, measurable progress and recorded effort; reviewable plans and evidence-based adjustments
Insights	Explain where time went and what creates inbox noise	Calendar/session and email statistics, with coverage and freshness; descriptive reports and proposed fixes

These are logical responsibilities, not a requirement for five autonomous agents or separate deployed services. They share a single context layer and durable records. Deterministic code owns checks, ranking rules, dates and capacity calculations; bounded AI steps extract or classify new text and write drafts, summaries and coaching.

Product model

A season is an optional longer planning container. A goal defines an outcome and success measure. A plan connects milestones, tasks, time allocations and follow-up loops to that outcome. A support area holds ongoing responsibilities without an artificial goal.

A commitment records something owed, the responsible person, who expects it, its source, due date if known, and resolution evidence. Tasks are actions that fulfil commitments; they are not duplicate commitment records. Independent personal tasks remain supported. A calendar block reserves time; a focus session records effort. Task completion, commitment resolution and goal progress are separate decisions and evidence.

Planning horizons are one day, one week, two weeks, four weeks, six weeks and twelve weeks. Longer plans do not reserve every future hour by default. Daily coaching preserves accepted blocks and proposes changes only when constraints or user priorities change. Users can mark a block as fixed; conflicting work is surfaced for a decision rather than silently moving it.

Loops and authority

A loop is an instance of a fixed, versioned, typed template with selected declarative parameters. Its specification defines trigger types, scope filters, deterministic checks, permitted action types, an authority ceiling, context recipe, model tiers, run/token caps and a stopping condition. Chat fills in that specification and presents what runs, when and why for confirmation. Parameters can lower authority but cannot exceed the template ceiling. Unsupported requests receive an explanation and the nearest supported option; the miss is recorded for product review.

The six starter template families are Daily brief, Daily check-out, Goal check-in, Overdue and slipping work, Focus rescue, and Weekly review. Friday wrap and week-ahead prep are modes of Weekly review. Daily coach sync is the interactive planning conversation that uses the brief; it is not a seventh ambiguous background template. Specialized assistant capabilities, including reply debt, waiting-on-them, VIP watch, meeting follow-through and meeting dossiers, run as built-in checks/output profiles inside these six templates. They are not additional top-level templates; see Automation for their registered mapping.

Free-text preferences may personalize output wording, length and order. They cannot alter checks, triggers, permissions or executable actions. Tasks, drafts and calendar actions inferred by the assistant are presented for approval; reading, deterministic analysis and briefing generation within granted access can run in the background. Correspondence and follow-up messages require exact per-message approval. Activating a notification loop authorizes its scheduled push/email delivery within configured limits. Every assistant-proposed schedule adjustment requires approval.

First-release boundary and build order

The first public release must deliver a full executive-assistant experience across the five engines, including commitments from email, supplied meeting material and manual entry. Calendar/goal/focus-only milestones are internal increments, not the released product.

Build Google Calendar reading, synchronization and approved publication with goals and tasks first; add meeting ingestion and follow-through next; then email ingestion and assistant workflows. Complete and verify the integrated assistant before first-release acceptance. Gmail is the proposed initial mail provider; supplied notes/transcripts are the first meeting input. Confirm import channel details before implementation; additional providers remain later work.

The web app comes first. It includes the calendar workspace, coaching, assistant workflows, in-app floating timer, optional lo-fi audio and Chrome website blocking. Desktop companions for macOS, Windows and Linux follow, including a system-floating timer where supported. Later system-wide website blocking starts with one selected desktop operating system. Native application blocking and cross-device blocking are excluded promises. The device owns session expiry and unblock recovery.

Background operation and cost

Server-side loops run even when the UI is closed. Each template defines a catch-up policy for missed runs. In-app results are the record of truth; web push and email fallback follow delivery preferences, permission and failure rules. Notification tiers are quiet, scheduled, nudge and urgent, with quiet hours and a daily limit.

Sources synchronize once into the shared store. Rolling summaries and content hashes support incremental processing; routines query the store instead of independently rereading providers. Context packs include bounded structured facts and source references. Usage is recorded per run and purpose, with reserved daily-brief capacity, per-run/per-loop caps and a global daily stop mechanism. When AI budget is constrained, fall back to clearly dated cached summaries or a deterministic digest and visibly pause lower-priority AI work.

Success measures

Prioritize commitment resolution, useful follow-up acceptance, timely reply handling, accurate and ready-before-open briefs, and reduced missed obligations. Also track time to first approved plan, goal outcome check-ins, session completion, proposal edit/rejection, notification dismissals and AI cost per purpose. Separate calendar reservations, recorded effort, task completion and measured outcomes. Set pilot targets after baseline usage; collect only consented telemetry.

Boundaries

Keep the interface minimal: calendar as the main workspace; contextual panels and drawers instead of a page for every routine. Use original Kairos copy and identity. Avoid invented progress, unexplained scores and automatic rescheduling. Broader connectors, native audio capture and AI handoff require separate capability and consent validation; they are not prerequisites for importing meeting notes into the first release.

Requirements & acceptance

All F requirements below belong to the first full executive-assistant web release. Build increments are internal validation stages, not progressively weakened release promises. Later-platform requirements are labeled L. Proposed numerical caps are intentionally not invented; they must be configured before any AI pilot and tuned from measured usage.

Product requirements

ID	Capability	Acceptance criteria
F01	Goal and season planning	User selects day/1/2/4/6/12-week horizon; defines success, path, capacity, constraints, follow-up and authority. Adaptive questions and manual fallback work; draft can resume.
F02	Plans and trajectory	Proposed milestones, tasks and time allocations show assumptions and capacity gaps. Measured outcomes, completed actions and effort remain distinct; insufficient data never produces invented progress or velocity.
F03	Google Calendar foundation	Selected permitted calendars supply freshness-aware availability/details. Approved publication and changes reconcile using source versions and stable operation IDs. Expired sync cursors trigger controlled recovery.
F04	Calendar-first interaction	Today/Plan share goals/tasks left, weekly/day calendar center and collapsible coach right. Drag/drop, resizing and equivalent keyboard/forms work; source, recurrence and external effects are visible.
F05	Daily briefing and coaching	Server prepares morning brief before app opening. Interactive coach preserves accepted blocks, asks about changed priorities/capacity and proposes only justified differences. User may skip sync.
F06	Exact approvals	Inferred task creation, saving a draft to an external mailbox, message sending and assistant calendar changes have versioned reviewable effects. Reports can render automatically. Direct user commands are authorization only for their explicit scoped effects.
F07	Stable schedule	Keep this time fixed prevents automatic placement/move/removal suggestions for that block. Conflicts are still surfaced; user can explicitly unlock or request a scoped change. No approval silently clears the lock.
F08	Commitment ledger	Email, supplied meeting notes and manual entry create sourced candidate commitments. Accepted commitments link fulfillment tasks; duplicate/ambiguous evidence is reviewed. Promise tracker, reply debt, waiting, meeting follow-through and VIP watch are available.
F09	Resolution evidence	Task completion is explicit; an incoming reply is not automatically resolution. Waiting items resolve only on clear evidence or confirmation. Goal completion uses its success criterion; evidence and reopening history remain inspectable.
F10	Calendar guardian	Checks cover conflicts, missing links, overload, merged Focus rescue, freed time and travel buffers. Outputs are findings/proposals. No physical travel estimate is invented without sufficient location/transport data.

ID	Capability	Acceptance criteria
F11	Meeting assistance	User-supplied notes/transcripts support decisions, candidate commitments and follow-up proposals; dossiers use available history with provenance. Native recording is not required for this release. Unknown owners/dates remain unknown.
F12	Email assistance	Gmail, the proposed launch mail baseline, supports bounded initial history, incremental ingestion, triage, reply debt, waiting evidence and draft proposals. Sending requires exact approval; permissions and provider acceptance are release gates. Outlook mail is later.
F13	Fixed workflow setup	One short chat fills a supported template's typed parameters. Preview states what/when/why/sources/authority/stop condition; confirmation activates. Unknown requests offer nearest supported configuration and record a minimal redacted miss.
F14	Predictable workflow logic	Six template families expose registered typed check/output profiles. Parameters cannot exceed template authority; free-text style preferences affect prose only. Template version changes never silently expand actions or sources.
F15	Server operation	Runs occur with the browser closed. Every template defines catch-up, expiry and deduplication. Late reports remain useful; obsolete urgent notices are suppressed. Paused/stopped loops cannot emit new deliveries.
F16	Managed delivery	In-app feed is durable. Web push is permission-aware; email fallback is opt-in and avoids duplicate sensitive content. Quiet/scheduled/nudge/urgent tiers obey quiet hours and configured daily limits.
F17	Focus and audio	In-app movable/minimizable timer and focus workstation support pause/resume/end, reload recovery, optional licensed lo-fi/ambient audio and Chrome website blocking. No native-app or cross-device protection is claimed.
F18	Focus evidence	Recorded active time excludes pauses; sleep/restart handling follows a documented policy. Completion records accomplishment, task status and optional outcome/self-rating separately. Device-local expiry and escape work without server access.
F19	Weekly review and insights	Friday wrap/week-ahead share review evidence and differ by mode. Time-use and inbox-noise statistics label coverage and distinguish scheduled, logged and inferred time. Carry-forward is a decision, not silent rollover.
F20	Shared context and costs	Sources ingest centrally; recipes build bounded source-linked packs. Deterministic checks precede AI. All runs record actual/estimated/unknown usage distinctly; purpose budgets protect briefing capacity.
F21	Proposal hygiene	Related changes bundle with dependencies; obsolete alternatives are superseded. Dismissals suppress equivalent proposals until relevant facts change. Partial batch results and conflicts are explicit.

Nonfunctional requirements

ID	Requirement	Acceptance criteria
N01	Access and consent	Ownership checks cover every entity. Source-reading, AI-sharing, notifications, publication and device-control permissions are explicit and revocable. Imported text never grants authority.

ID	Requirement	Acceptance criteria
N02	Reliable state	Retry-safe ingestion, jobs, approvals, notifications and provider effects have stable keys. Unknown remote outcomes reconcile before retry. Record partial success; do not promise atomic writes across providers.
N03	Time correctness	IANA timezone, all-day bounds, DST, travel timezone changes, quiet hours and catch-up are defined/tested. Time passing invalidates time-sensitive cached results.
N04	Accessibility and layout	Keyboard alternatives, visible focus, labels, reduced motion and text zoom work. Narrow screens use day/agenda plus drawers. Status is not conveyed by color alone.
N05	Honest degradation	Manual planning works without AI; cached content has as-of/coverage labels. Low budgets fall back to fresh deterministic facts; new unsupported claims are not generated from stale summaries.
N06	Privacy lifecycle	Each source defines retention, encryption, model-sharing boundary, export, disconnect and deletion for raw data, derived summaries and audit records before release.
N07	Performance and evaluation	Agree representative data/device budgets before UI acceptance. Evaluate extraction, duplicate detection, commitment resolution and briefing grounding with labeled scenarios.
N08	Budget enforcement	Configure conservative caps before a pilot; tune after measured use. Atomic reservations prevent concurrent overspend. Global kill switch blocks AI, not local unblock or manual access.

Later platform scope

ID	Scope	Gate
L01	Desktop timer companions for macOS, Windows and Linux	After web release; platform-specific signed packaging, permissions, recovery, updates and acceptance. No app-blocking implication.
L02	System-wide website blocking on one selected desktop OS	Later opt-in hosts helper or reviewed equivalent; reversible, privilege-bounded, locally expiring; disclose coverage limitations. No native-app/cross-device blocking.
L03	Additional providers, native capture and AI handoff	Separate validated connector/capture contracts and user authority; not required for first web release.

Approval interpretation

A report/brief, deterministic finding or in-app preview is not an external mutation. Candidate extraction can persist suggestions in the ledger without accepting them as obligations. A generated reply can be displayed automatically as preview; approving creation of an external mailbox draft does not approve sending it. Direct explicit task creation needs no redundant generic approval. Every assistant schedule adjustment still needs review, with a single batch approval permitted only when all effects and destinations are listed.

End-to-end workflows

1. Connect and clarify

> 1. Choose plan a goal or organize existing work; select a day/1/2/4/6/12-week horizon.
>
> 2. Connect Google Calendar, then Gmail when ready, with separate source/AI/write consent. Enter availability manually if not connected.
>
> 3. Define success, path, capacity, constraints, follow-up and authority. Follow-up questions adapt to missing information; save/resume is available.
>
> 4. Show bounded import scope, progress, partial coverage and source freshness. Do not imply older history was checked.
>
> 5. Propose milestones, tasks, blocks and workflows; expose assumptions/capacity gaps. Review exact effects and destinations, accept/edit/reject.
>
> 6. Open the calendar-first workspace. Unapproved items remain suggestions; publication is approved explicitly, possibly within the same enumerated review.

2. Activate a fixed workflow

Choose one of six template families, answer a short chat, inspect the typed form preview and confirm. Chat selects only registered check/output profiles and allowed parameter values; it cannot invent new checks, code, tools or permissions. Style instructions go only to writing. Unsupported requests offer the nearest template and record a redacted catalog miss. Preview shows next run, source coverage, authority ceiling and stop condition. Changes to behavior are reviewed as a new version.

3. Morning brief and daily coach sync

The server creates a Daily brief before the app is opened. Today shows it in the coach panel beside the calendar. The interactive daily coach sync uses that report; it is not a separate automation template. Coach asks about changed capacity, energy and core priorities. Preserve the accepted weekly plan and fixed blocks. Propose only justified additions/differences, with reasons. Review and approve individual changes or an exact batch, then start focus. Skip sync without losing the existing plan.

4. Capture and resolve commitments

Email, meeting notes or manual input supplies evidence. Extract candidate commitments and match people/threads with provenance. User accepts, edits, dismisses or merges candidates before inferred work enters the accepted ledger. A commitment records what is owed; tasks describe actions that fulfill it. One task can contribute to an obligation without proving that obligation resolved.

Resolution rules:

- Task: explicit user completion or a separately approved, verifiable completion action.

- Waiting on them: a response creates evidence; mark resolved only when deterministic evidence is conclusive under the configured rule, otherwise request confirmation. Acknowledgments and partial answers do not close it.
- Promise/meeting action: completion evidence must satisfy the promised deliverable, with owner confirmation when uncertain.
- Goal: record the success measure; effort and task count are not substitutes.
- Deferred work: capture reason, resume date or review trigger; pause follow-ups until that trigger. Reopen when evidence invalidates an earlier resolution and show why.

5. Guardian and approval review

A source event or time-based check identifies conflict, freed capacity, missing link, travel need or threatened focus. Focus rescue merges preventive Focus guard and reactive Focus defense. Calendar sentinel remains a set of guardian checks, not a duplicate routine.

Produce a bundled proposal with before/after, dependency ordering, source versions and exact destinations. Fixed blocks remain protected; surface a decision rather than moving them. Preserve the currently observed external calendar: approval governs Kairos's corrective action, not whether a third-party edit already happened.

Revalidate on approval. Supersede old alternatives; changed facts require renewed review. Remember dismissal signatures (reason, affected items, meaningful constraints) until those facts change. Applying one change cannot silently apply a dependent unapproved change. Report partial results and reconcile unknown external writes before retrying.

6. Focus

Start a task/block, confirm objective/duration and configured Chrome website rules. Optional audio starts through user playback controls. Dedicated view and floating in-app timer show active time, pause/resume/end and actual browser protection status. Pause releases session blocking by default; resuming restores permitted rules. Emergency unblock and local expiry remain available without network.

On finish, record accomplishment, explicit task status, remaining work and optional outcome/self-rating. Timer expiry never completes a task. Sleep, restart and abandoned sessions need a declared recovery policy before implementation; do not claim idle/sleep time as focused work. A continuation block is a new proposal. Later desktop timer companions do not imply native-app or cross-device blocking.

7. Daily check-out and weekly review

Daily check-out reconciles completed, unfinished, waiting and blocked items; `/done` opens this session. Updates explicitly entered by the user apply as entered; inferred new tasks, mailbox drafts and calendar changes remain proposals. Missing the check-out does not mark all unfinished work failed or rolled forward.

Weekly review has Friday-wrap and week-ahead modes over one evidence snapshot/review record. Friday summarizes outcomes and open decisions; week-ahead proposes next week. Deduplicate overlapping reports; preserve differences when constraints change. Compare scheduled time, recorded active time and coverage, not invented actual activity. User decides carry-forward, defer, drop or revise.

8. Closed-app operation and catch-up

Ingestion, checks, briefing and approved notifications run server-side while the browser is closed. The device independently ends website rules. Feed stores report and decision history; push/email provide permitted delivery, not guaranteed proof the user read it. On reconnect, show latest useful outputs and partial-source warnings rather than replaying every old notification. Detailed catch-up and delivery rules are in Automation.

State contracts

Plan: draft → proposed → active → completed/archived; prior accepted revision remains authoritative until changed.

Commitment: candidate → accepted → active/waiting/deferred → fulfilled/cancelled. Candidate can be dismissed; deferred needs a resume trigger. Tasks have independent todo/in-progress/done/cancelled state and scheduling links. Goal outcome state is separate.

Proposal: pending → applying → applied/partially-applied/failed; pending can be dismissed/superseded/expired/conflicted. Retries retain operation identity. Changed effects create a new review version.

Loop: draft → active ↔ paused → stopped. Run: queued/running → success/partial/failed/skipped/cancelled. Check pause/version before delivery; cancelled work cannot emit newly generated actions.

Focus: idle → running ↔ paused → ended; local end/escape is authoritative for device rule release. Reconciled timestamps must not duplicate duration or extend expired protection.

Interface architecture

Clean Tech visual system

White and light-gray surfaces, deep charcoal typography, ample white space and subtle borders. One vibrant deep-purple accent (#6438E8) supports primary buttons, links, selection and focus. A pale tint of the same accent is permitted. Neutral labels/icons/outline patterns distinguish categories, sources and statuses; no multi-color pastel category palette. Original Kairos text and assets follow the supplied reference's hierarchy and simplicity.

Six views, one calendar-first entry

View	Role
Today	Default entry: calendar-first shell anchored to today, with morning brief/daily coaching and urgent decisions. No metric-heavy replacement dashboard.
Plan	Same shell for week-ahead work, full calendar, tasks, proposals and capacity. Preserve date/selection when switching context.
Commitments	Accepted/candidate obligations, replies owed, waiting, promises and linked tasks; inspect evidence and resolution.
Goals	Seasons, goal setup, milestones, plans and trajectory; task scheduling returns to the same calendar.
Review	Daily check-out, Friday wrap/week-ahead modes, outcome and coverage-aware effort insights.
Assistant	Template setup through short chat, workflow settings/history, preferences, usage and connected sources.

Workspace composition

Coach closed: left goals/tasks/support panel approximately 26%, calendar approximately 74%, with practical min/max widths. Coach open: protect a minimum usable calendar and collapse the left panel or show the coach as an overlay at laptop widths. At an illustrative 1440 px viewport: 24 px outside margins each, 272 px sidebar, 760 px calendar, 312 px coach and two 24 px gutters sum to 1440 px. This is a prototype example, not a universal breakpoint rule. An overlay is preferable when a seven-day grid becomes unreadable.

Left panel has quick capture, expandable goals/tasks, support areas, filters and sources. Center has date navigation, day/week switch, all-day row, hourly grid, now marker, source freshness and proposal review strip. Right panel holds the brief, interactive daily coaching, follow-ups, approvals and chat. A compact primary navigation avoids adding another oversized sidebar.

Interaction rules

Drag tasks to schedule; offer equivalent keyboard/form controls. Inspect event details in a popover/drawer. Direct edits expose local/external effects. Assistant additions use dashed outlines, moves show before/after, and proposed removals remain marked until approved. Keep this time fixed is visible on blocks and respected by planning. Third-party changes display as observed facts; corrective changes require approval.

Proposals bundle related changes, show dependencies and destinations, suppress equivalent dismissed suggestions and replace obsolete versions. A concise card expands for evidence and detailed effects. Approval is specific to a version; chat is not a general grant of authority.

Focus

Dedicated task-centered view with optional audio, in-app timer and Chrome website-rule status. No native-app or cross-device protection indicators. Device owns expiry and escape. Later desktop companions supply timer access on macOS/Windows/Linux; later system website blocking starts with one supported OS. All device capabilities are labeled accurately.

Responsive and accessible

Desktop prioritizes the calendar; medium screens use drawers; phone uses day/agenda and task/coach sheets. Preserve selected date and task. Provide keyboard alternatives, focus restoration, text zoom and reduced motion. Status uses words/icons as well as styling. Loading, empty, partial, stale, denied, offline, failed and budget-limited states are explicit. See screen specifications for components and step-by-step flows.

Screen design & components

Design version: 2.1 / 31 October 2026. Clean Tech is the single visual direction. This revision includes full assistant workflows, approval hygiene and web-first website-only focus controls. Dimensions below are proposed design targets to validate in a prototype.

1. Visual system

White #FFFFFF surfaces; light gray #F6F7F9 canvas; charcoal #17191F primary text; slate #596170 secondary text; subtle #E2E5EB borders. Deep purple #6438E8 is the sole vibrant accent for primary actions, selected controls, links and focus indicators. Pale purple #F0EBFF is a tint of that same accent. Use neutral labels and outlines for event categories, not additional hues. Success, warning and error states use explicit text and icons so meaning never depends on color.

Use a clean sans-serif typeface, 32-40 px major headings, 20-24 px screen headings, 14-16 px body and 12-13 px supporting metadata. Minimum sizes and contrast require device verification. Use an 8 px spacing rhythm with 4 px fine adjustments; 16-24 px panel padding, 8 px event corners and 12 px panel corners. Borders supply separation; shadows are reserved for floating surfaces. One primary action per decision context.

2. Screen S01 / Welcome and context

Purpose: establish the user's intent without forcing a long setup interview.

Layout: slim brand header, centered content column, progress indicator and a visible Save and exit action. Present two starting choices: plan a goal, or organize existing work. Explain what the assistant will prepare. Google Calendar and Gmail connection follow with separate permissions; manual availability and supplied meeting notes remain available. Declining connections shows reduced assistant coverage honestly.

Steps: choose intent; select planning horizon; describe the desired outcome; connect sources or continue manually. Next advances only after required information is supplied. Retain entries when navigating back. Connection denial explains what remains usable.

Components: intent cards, horizon selector, text field, source connector, progress stepper, inline validation and primary Continue button.

3. Screen S02 / Goal clarification

Purpose: turn a broad ambition into an actionable brief.

Layout: structured goal form with a compact conversational follow-up area. Show known answers as editable fields, not buried only in a transcript. Ask for success measure, current position, deadline, available time, priority, constraints and dependencies. Distinguish required answers from optional coaching context.

Steps: capture goal; ask the most relevant unanswered question; summarize assumptions; let user correct; generate the proposed plan. Do not require every possible question for a simple one-day task. AI unavailable state offers a manual goal and task form.

Components: outcome fields, date picker, capacity input, editable answer chips, question card, draft-save status and Generate plan action.

4. Screen S03 / Plan review and activation

Purpose: make the assistant's proposal concrete and reviewable.

Layout: outcome summary above milestone/task list; calendar preview alongside on wide screens; proposed loops below. Each item shows purpose, estimate, date or frequency, and assumptions. A capacity summary explains unused time or overcommitment.

Steps: inspect the plan; expand supporting reasons; edit tasks or timing; choose proposed items; review listed effects; approve. External publication names the calendar destination and effects explicitly. Keep local acceptance and remote publication distinct unless a single approval explicitly includes both.

No capacity: offer reduce scope, extend deadline or change availability. Rejected suggestions leave accepted work intact. Changed availability invalidates affected suggestions and asks for renewed review.

Components: milestone list, task editor, capacity summary, calendar preview, loop settings card and approval summary.

5. Screen S04 / Calendar home

Purpose: manage priorities and real-world time in one workspace.

Desktop: goals/tasks on the left, calendar in the center and collapsible coach on the right. With coach closed, use approximately 26% / 74%; constrain sidebar to a usable range. At 1440 px, an illustrative open layout is 272 px sidebar, 760 px calendar, 312 px coach, two 24 px gutters and 24 px outer margins on each side, totaling 1440 px. Collapse the sidebar or overlay the coach when the calendar becomes too narrow. This is a prototype target, not a fixed rule.

Left: quick capture; expandable season goals with tasks; support areas; calendar/source filters. Center: date navigation, Day/Week, all-day row, hourly grid, current-time marker, source freshness and proposal strip. Right: today's sync, pending decisions, follow-ups and contextual chat.

Steps: expand a goal; drag a task into time or use Schedule; inspect event; edit through popover; start focus. Dragging assistant-proposed changes does not bypass proposal review. Direct user edits identify any external save effects. Clicking a blank slot creates an explicit task/event/block choice.

Five- or seven-day display is a user preference; default seven-day week follows the reference. Work-hour display can be expanded. Source labels distinguish Google events, local events and goal blocks without a category rainbow. Keep this time fixed is visible on blocks; routine proposals respect it. Show third-party calendar edits as observed facts immediately, then offer corrective proposals if needed.

6. Screen S05 / Daily coach sync

Purpose: agree on today's core work before focus.

Panel structure: concise daily brief; today's constraints; proposed top priorities; proposed time blocks; pending decisions; chat input. Start with actual goals and available time. Ask about changed capacity and blockers. Explain priority tradeoffs briefly, with evidence links available on demand.

Steps: review the precomputed server brief; correct priorities; tell coach today's capacity; preserve accepted blocks and fixed times; inspect only justified proposed differences; approve individual changes or the enumerated batch; start selected task. Skipping sync preserves the existing plan. The assistant does not interpret a conversational acknowledgment as blanket authorization for unspecified changes.

Components: brief card, priority list, capacity prompt, proposal cards, evidence disclosure and Start focus action.

7. Screen S06 / Proposal review

Purpose: approve every assistant-proposed schedule adjustment.

Layout: compact summary strip in calendar; detailed drawer for review. Each change shows action type, previous state, proposed state, reason, affected task, source and destination. Additions use a dashed outline; moves show old/new positions with labels; removals remain visible until approved.

Steps: select a proposal; inspect before/after; edit or dismiss; approve; revalidate current calendar; apply; report local and external outcomes separately. A stale proposal shows the conflict and offers a refreshed version. Failure must preserve accepted state or explain any partial effects.

Components: change card, before/after comparison, dependencies, selection checkbox, destination summary, Apply selected, conflict banner and retry/reconcile action. Repeated approval is idempotent. Bundle related changes, supersede obsolete alternatives and remember equivalent dismissals until relevant facts change. Task adoption and external mailbox draft creation are reviewable effects; an internal generated preview is not an accepted task or sent message.

8. Screen S07 / Focus workstation and floating timer

Purpose: protect attention while keeping essential controls accessible.

Workstation: task title and intended outcome, timer, pause/end controls, Chrome website-blocking status and optional audio. Do not show native-app or cross-device protection. Hide nonessential calendar and chat content during focus, while allowing intentional return to planning. Starting a session uses only configured supported blocking rules.

Floating timer: compact neutral surface with purple active indicator, task label, elapsed/remaining time and pause/end. Drag handle, minimize and keyboard move controls; restore a safe position on viewport change. The first scope is in-app floating. Desktop companions with floating timers follow the web release on macOS, Windows and Linux; one-OS system website blocking is a separate later feature, not native-app blocking.

Steps: select task; confirm duration/objective; start; verify supported Chrome/profile website-rule status; optionally play audio; pause/resume; finish. Proposed pause behavior releases blocking while paused. Show unavailable controls honestly. Emergency unblock works even when the server cannot be reached; timed rules expire independently.

Audio controls: play/pause, volume, track/source and mute. No automatic sound before a user gesture. Audio failure leaves the timer usable. Licensing/source selection remains a delivery decision.

9. Screen S08 / Session completion

Purpose: record evidence and choose a useful next step.

Layout: short completion sheet with actual recorded duration; accomplishment prompt; explicit task status; optional outcome measure and optional self-rating. Do not infer task completion from timer expiry. Do not show a precise focus score without defined observation inputs and validation.

Steps: describe result; mark task complete or unfinished; record measurable progress if applicable; capture blocker; save. Offer a proposed continuation session or return to calendar. Saving session evidence does not approve a new calendar adjustment.

Components: duration summary, status control, note field, measure input, optional rating and next-step proposal.

10. Screen S09 / Follow-ups and loop settings

Purpose: let the user understand and control ongoing assistant work.

Follow-ups are coach-panel cards grouped by urgency and routine. Expand to inspect evidence, task owner, due date, draft and proposed action. Email sending and external changes require explicit authorization. Dismiss hides a finding; snooze changes reminder time; neither means the underlying commitment is complete.

Assistant exposes six template families configured through one short chat and a typed confirmation preview. Loop detail contains pinned template/profile version, purpose, goal/area scope, schedule, timezone, permitted sources, deterministic criteria, authority ceiling, stop/catch-up condition, quiet hours, notification destination, next run and last result. Chat cannot exceed the registered schema; style preferences affect writing only. Pause stops future runs; stop terminates the loop. Keep run history readable and errors actionable. Partial results name unavailable sources without exposing credentials or raw internal errors.

11. Screen S10 / Weekly review

Purpose: learn from the week and prepare a realistic next one.

Layout: outcome progress, completed work, scheduled versus recorded effort, repeatedly moved work, unresolved commitments and short reflection prompts. Focus time and goal outcomes remain separate. Place proposed next-week adjustments beside a calendar preview rather than burying them in prose.

Steps: inspect evidence; correct missing progress; decide carry-forward/drop/defer; discuss constraints; generate next-week proposal; approve exact changes. Defer changes a commitment deliberately; it is not silent rollover.

Components: outcome summary, effort comparison, carry-forward list, reflection form and next-week proposal review.

12. Component contract inventory

Component	Main behavior	Required states
Goal tree	Expand tasks; select scope; add goal	Empty, expanded, filtered, archived
Task row	Complete, inspect, schedule, drag	Candidate, todo, in-progress, done, cancelled
Calendar block	Inspect, move, resize, start focus	Local, external, proposed, stale, read-only
Proposal card	Explain, edit, accept, dismiss	Pending, superseded, conflicted, partial, applied, failed
Coach message	Clarify with source context	Streaming, cancelled, unavailable, retry
Floating timer	Move, minimize, pause, end	Running, paused, reconnecting, ended
Browser rule status	Report supported Chrome website rules	Active, paused, denied, offline, expired

Component	Main behavior	Required states
Loop card	Configure, pause, inspect results	Active, paused, failed, partial, stopped
Event popover	Edit relevant details	Read-only, editable, recurrence choice, conflict
Source connector	Connect, scope, disconnect	Unconfigured, syncing, current, stale, revoked

13. Responsive and accessibility behavior

Desktop keeps the calendar dominant. Tablet uses a coach drawer over the workspace. Phone defaults to day/agenda view with separate task and coach sheets; no squeezed seven-column grid. Preserve selected date/task when changing surfaces.

Support keyboard scheduling and resizing alternatives, visible focus, descriptive labels, meaningful dialog focus restoration, screen-reader announcements for approval and timer changes, reduced motion, text zoom and comfortable touch targets. Use text and icons for all statuses. Do not announce timer ticks continuously to assistive technology. Validate accessibility and measured contrast in the prototype and implementation.

14. Design validation walkthrough

Walk a user through: define a six-week goal; answer clarification; revise a proposed plan; activate a loop; run daily coach sync; schedule a task without dragging; approve a move; start focus; navigate with the timer; pause/unblock; finish without completing the task; handle stale Google data; review the week. Confirm comprehension of proposed versus accepted changes and which devices are actually protected.

This document specifies screen behavior, not a claim that the screens or all integrations are implemented. The web/desktop sequence and no-app/no-cross-device scope are fixed. Exact supported versions, first website-helper OS, audio licensing, session recovery and pilot targets remain open decisions.

15. Supporting views S11-S14

Today and Plan share S04: Today anchors the calendar and brief to the current day; Plan emphasizes week-ahead capacity and proposed changes. The shell remains calendar-first rather than introducing a competing home dashboard.

View	Content and main steps	Components and critical states
Commitments / S11	Filter candidates, mine, waiting, reply debt and promises; inspect evidence; accept/merge; schedule fulfillment tasks; confirm resolution or defer with resume trigger	Ledger row, owner/due/source labels, evidence drawer, task links, merge review, partial-answer and reopened states
Goals / S12	Define season/outcome; inspect plan, tasks and recorded progress; revise priorities; return to calendar for time decisions	Goal detail, outcome measure, capacity, milestone list, fixed-block links, insufficient-evidence trajectory

View	Content and main steps	Components and critical states
Review / S13	Daily check-out, Friday wrap/week-ahead modes, scheduled-vs-logged time and inbox-noise coverage	Shared review record, carry-forward decisions, coverage labels, fresh deterministic fallback
Assistant / S14	Pick template; short setup chat; confirm preview; manage sources/preferences/usage; inspect history; pause/stop	Template picker, typed configuration, authority diff, unsupported-request feedback, usage meter, degraded/paused states

16. Background and budget UX

Feed items show produced-at/as-of, source coverage, notification tier and action status. Push permission denial offers an in-app path and opt-in email fallback; no duplicate fallback merely because an alert was not read. Missed runs coalesce into the latest useful finding. Respect quiet hours and daily caps while keeping pending decisions visible.

Usage meter shows scheduled loops, extraction, drafting and chat; identify reserved daily-brief capacity and one near-80-percent alert. Budget degradation is visible: full output, cached summaries with current deterministic checks, current deterministic digest, or paused low-priority AI. Never interpret unprocessed new mail as no reply debt.

17. Full assistant acceptance walkthrough

With the browser closed, a new permitted email and supplied meeting note are ingested. Candidate promises and waiting items are linked to evidence. The morning brief is generated before opening; permitted push/email follows delivery rules. Open Today, inspect the brief, accept an inferred task, reject a calendar move and verify it does not reappear unchanged tomorrow. Resolve a partial reply without prematurely closing the request; inspect a meeting dossier, approve a follow-up draft/send separately, run focus with Chrome website rules, check out and review the week. Repeat with revoked permissions, stale sources, expired cursor, low AI budget and a fixed block collision.

Features & integrations

Version 2.1 / 31 October 2026. Planned release scope; this catalog does not certify existing implementation.

Release contract

The first released web product is an executive assistant across the commitment ledger, calendar guardian, briefing composer, goal engine and insights. Calendar/goals/tasks, meetings, then email is implementation order. Internal increments do not redefine the public release as a planner-only MVP.

Five-engine feature inventory

Cadences below are proposed defaults from the requested workflow inventory, configurable within each template's allowed values and the user's timezone. All assistant actions remain reviewable proposals; detected facts and prepared reports may appear automatically within granted access.

Engine	Fixed check or output profile	Default cadence/trigger	Result
Commitment ledger	Promise tracker	Daily and relevant source changes	Unresolved promises, owners, due dates and evidence
Commitment ledger	Waiting on them	Twice weekly	Outstanding requests and nudge drafts
Commitment ledger	Reply debt	Daily	Ranked unanswered requests with reply drafts
Commitment ledger	Meeting follow-through	After supplied notes arrive; evening recap	Proposed commitments/tasks and follow-up drafts
Commitment ledger	VIP watch	Configured threshold checks	User-selected VIP attention findings and draft replies
Calendar guardian	Calendar sentinel	Daily and relevant calendar changes	Double bookings, missing links and overloaded days
Calendar guardian	Focus rescue	Daily/weekly reservation check and relevant changes	Unified focus protection/recovery proposals; replaces separate focus guard/defense routines
Calendar guardian	Freed time	Meeting cancellation or shortened event	Suggested use of the newly available slot
Calendar guardian	Travel buffer	Relevant in-person event change	Proposed travel reservation, with route/duration assumptions visible
Briefing composer	Daily brief	Morning before configured delivery time	Priorities, obligations, meeting prep and pending decisions

Engine	Fixed check or output profile	Default cadence/trigger	Result
Briefing composer	Meeting dossier	Before external meeting	Available relevant history, open threads and promises
Briefing composer	Friday wrap	Weekly review: wrap mode	Accomplishments, unresolved work and carry-forward decisions
Briefing composer	Week-ahead prep	Weekly review: preparation mode	Next week's decisions, constraints and proposed plan
Goal engine	Goal check-in and slipping-work checks	User-selected schedule or supported threshold	Outcome check-in, stalled work and specific proposed fixes
Goal engine	Season/GPS-style planning and weekly plan	On explicit planning request	Clarified goals, milestones, tasks, capacity-aware blocks and loops
Goal engine	Trajectory and velocity	After relevant recorded progress	Evidence-based trend with coverage and uncertainty
Insights	Where your week went	Weekly	Scheduled versus recorded effort, categories and trends
Insights	Inbox noise	Every two weeks	Email volume patterns and proposed rules

Travel buffers must work with user-entered travel time or an explicitly chosen routing source; the assistant must not invent travel duration. Meeting and email coverage limitations remain visible. Relationship-cooling analysis is later scope, outside the stated first-release Insights inventory.

Six starter template families

Template family	Scope	Purpose and supported modes
Daily brief	Global	Prepare the morning brief from shared facts and pending decisions
Daily check-out	Global	Evening reconciliation and `/done`; capture outcomes and unfinished work
Goal check-in	Per goal; ledger checks also allow global/area scope	Review progress and goal-specific next actions; promise/waiting checks use explicit scope exceptions
Overdue and slipping work	Global or per goal	Detect overdue commitments/tasks and repeatedly postponed work
Focus rescue	Global	Calendar-guardian checks and proposed protection/recovery
Weekly review	Global	Friday wrap and week-ahead prep modes using shared review data

Daily coach sync is the interactive conversation that turns brief/context into an approved day. It preserves accepted blocks, respects fixed-time choices and does not run a new full-week rescheduler every morning.

Confirmed workflow catalog

Confirmed 31 October 2026: reply debt, waiting-on-them, VIP watch, meeting follow-through and meeting dossiers run as built-in checks/output profiles inside the six templates. They are not additional top-level selectable templates. The five engines retain the full capability inventory above.

Each built-in check has registered trigger options, scope, deterministic criteria, outputs, context recipe and authority ceiling. Users tune and enable supported checks in the template setup rather than composing arbitrary automation. The registered check mapping supplies the design configuration; implementation must validate each schema and trigger. Family names do not restrict every check to the cadence implied by the family name.

Template and loop contract

Every template defines its ID/version, supported trigger kinds (schedule, source event, threshold), scope filters, deterministic check library, allowed output/action types (brief, notification, draft, proposal), authority ceiling, stop condition, catch-up policy, context recipe, model tier per AI step and token/step/tool-call caps. Each loop stores template ID/version and validated chosen values. Template upgrades cannot silently change user authority or behavior; changed behavior needs a reviewable migration/confirmation policy.

Chat chooses a supported template/profile and fills typed fields. A confirmation preview names what runs, when, which sources it reads, what it may prepare and when it stops. Values can reduce the authority ceiling, never raise it. Unsupported requests show the nearest supported option and log the miss. Writing preferences such as brief length and tone affect prose only, not trigger/check/action logic.

Approval and resolution

Bundle related changes into one review with dependencies. Replace obsolete proposals and remember dismissals until relevant facts change. Explicit approval covers the exact listed task, draft or calendar effects; correspondence never sends without exact per-message approval. Activated notifications use standing delivery authorization within their configured limits. External calendar publication identifies its destination. Accepted weekly blocks remain the baseline; new constraints or user priorities justify proposed adjustments. A fixed-time block is not moved by the assistant.

Commitments close using type-appropriate evidence. A personal task requires explicit completion; an incoming reply is evidence of a response, not proof that the request was resolved. Ambiguous meeting and email conclusions require confirmation. Goal completion follows its defined success measure. Reminders, elapsed time and sending a message do not automatically close the underlying obligation.

Integration and platform boundary

Integration/capability	First-release or later role	Verification boundary
Google Calendar	First-release foundation: read availability/details as permitted, synchronize, publish approved changes	Preserve separate consent, source capabilities, version conflicts and real-provider acceptance

Integration/capability	First-release or later role	Verification boundary
Gmail	First-release commitment extraction, reply/waiting/VIP checks, drafts and noise statistics	Proposed initial mail provider; least-privilege access, retention and provider verification required; additional providers later
Supplied meeting notes/transcripts	First-release meeting context, dossiers and follow-through	Select manual/import channel and retain attribution; do not imply automatic attendance/capture
Outlook calendar / Apple Calendar	Later calendar connectors	Validate supported protocol, read/write scope and conflict behavior before promising parity
Slack, Fathom, Granola	Later context/meeting connectors	Validate accessible source capabilities, consent and retention
Zoom / Meet / Teams / Webex capture	Later native audio capture	Local permission, supported OS/capture mechanism and recording notices; not a web-only promise
Chrome extension	First-release website blocking	Verify actual Chrome lifecycle and recovery behavior; no native-app blocking claim
Desktop companions	After web: macOS, Windows and Linux timer/session surface	Validate each OS; system-floating behavior is a device capability, not assumed browser parity
System-wide website helper	Later, one selected desktop OS first	User-approved bounded website rules, independent expiry and recovery; no application or cross-device blocking
External AI handoff	Later reviewed task/context transfer	Destination, context limits and completion/status contract required

The repository contains a legacy opt-in Linux helper that can suspend configured native executables. That existing code is outside the revised launch product scope; it is not evidence of promised macOS/Windows/Linux native-app blocking, and this documentation revision does not remove or alter it.

Background delivery and notifications

Loops run server-side while the UI is closed, with template-specific catch-up and deduplication. Store the result in the in-app feed first. Deliver through permitted web push, then email fallback according to configured preferences and delivery outcomes; do not send duplicate fallback messages merely because a user has not opened the app. Quiet, scheduled, nudge and urgent tiers obey quiet hours and daily limits, with any urgent exception explicitly configured.

Show next run, last run, partial-source warnings, missed-run handling and paused budget status. Delivery fallback rules are defined in Automation; exact catch-up windows and urgent-tier exceptions must be configured before activation. The device, not a server job, owns focus session expiry and unblock recovery.

Shared context and budget behavior

Adapters sync each source once into owned people, threads, meetings, commitments, goals and summary records. Routines query this store, not providers directly. Hashes and rolling summaries limit AI work to changed material; a bounded initial backfill limits first-connection processing. Unchanged facts can reuse summaries, but deadline/time-threshold checks still run as time advances.

Use code for overdue/waiting calculations, conflict detection, ranking and capacity math. Use a lower-cost validated AI step for extraction/classification and a stronger model only for approved writing/coaching paths. Context packs have source references and token ceilings; sub-steps share structured results rather than raw transcripts.

Record tokens and cost per run, loop and purpose. Reserve daily-brief capacity separately from chat; cap each run/loop and support a global daily stop mechanism. Provide a usage meter and one near-80-percent alert. Numeric quotas follow pilot observation; hard safety caps must still exist during the pilot. Degrade from full output to dated cached summaries, deterministic digest and explicitly paused low-priority AI loops. Never represent a cached output as newly analyzed evidence.

Templates & background operation

Fixed declarative templates

A template is a versioned typed specification, not generated code. It declares supported schedule/event/threshold triggers, scope filters, deterministic checks, output profiles, authority ceiling, stop conditions, context recipe, model tier per step, token cap and catch-up policy. A user's loop stores template ID/version and validated parameters. User values may reduce authority, never increase it. Unknown fields/actions fail validation.

Six user-facing template families: Daily brief, Daily check-out, Goal check-in, Overdue and slipping work, Focus rescue, Weekly review. Confirmed 31 October 2026: specialized assistant capabilities are built-in checks/output profiles inside these six templates, not additional top-level templates. All listed executive-assistant capabilities remain required in the first web release.

Chat fills the typed form, explains the preview and waits for activation confirmation. Unsupported requests identify the limitation, offer the closest supported configuration and record a redacted miss for product discovery. Free-text preferences such as short briefs or lead with money matters affect prose only; typed priorities determine checks/ranking. They cannot modify logic, sources, cadence, tools or action authority.

Registered built-in check mapping

Family	Fixed checks/output profiles	Trigger and context
Daily brief	Morning overview; VIP watch; meeting dossier	Morning schedule; VIP wait threshold; before-meeting event. Reads accepted priorities, open commitments, calendar and source-linked relationship/thread summaries.
Daily check-out	Evening reconcile; meeting follow-through; `/done`	Evening or meeting-note-ready event. Reads new meeting evidence, decisions, candidate commitments and today's tasks.
Goal check-in	Progress, trajectory, promise tracker, waiting-on-them	Scheduled/threshold checks scoped to a goal or selected commitment/area; global ledger profile has no mandatory goal. Reads outcome measures, tasks and relevant unresolved commitments.
Overdue and slipping work	Overdue tasks, repeated postponements, reply debt	Scheduled or threshold; global or goal scope. Reads due/wait dates, meaningful move history, urgency/VIP settings and unresolved replies.
Focus rescue	Focus guard + defense; sentinel checks; freed time; travel buffer	Calendar change or periodic check. Reads current availability, fixed blocks, capacity, travel fields and selected goals.

Family	Fixed checks/output profiles	Trigger and context
Weekly review	Friday wrap; week-ahead prep; time-use; inbox noise	Weekly modes; inbox-noise profile can use biweekly cadence. Reads shared review record and aggregate statistics.

Family names do not force one daily/weekly cadence on every registered profile. Profiles have their own schema-approved trigger choices. Users can enable multiple instances; deduplication prevents the same finding being reissued across profiles. The built-in-check architecture is confirmed. The mapping above is the design specification to validate through schema and workflow tests; adding a check does not create a seventh selectable template.

Authority and review

Automatic under granted permissions: ingestion, structured candidate extraction, deterministic checks, in-app draft previews, reports and permitted notifications. AI cannot create accepted obligations or execute external side effects on its own.

Proposed effects: accepted task creation from inference; external mailbox draft creation/update; calendar addition/move/removal; follow-up message sending. Review lists exact items/destinations and dependencies. Approving a draft is distinct from approving sending. Direct user commands authorize only the clearly requested scope. Every assistant schedule adjustment requires approval, including its own blocks.

Related effects share a batch review with itemized choices. A dependent effect cannot apply unless its prerequisites are approved/applied. Before applying, check owner, permissions, source/version/freshness, availability and locks. Superseded/expired or meaningfully changed effects require renewed review. Partial external success is reported; unknown outcomes reconcile before retry.

A dismissal records a meaningful finding signature, not just its generated wording. Suppress equivalent proposals until relevant source facts/constraints change. Time-based thresholds can constitute a declared change only when the template explains it; daily reruns alone are not a reason to revive dismissed suggestions.

Server runs and catch-up

Server scheduler and source-event processing run independently of the browser. Store a logical trigger key, input snapshot, template version, planned time, actual time, freshness, budget reservation and outcome. Lease/lock execution to prevent duplicate processing. Recheck pause/stop and template version before emitting outputs. Retry transient failures with bounded backoff; surface persistent failure and recoverable partial results.

Profile type	Missed-run rule	Obsolescence rule
Morning brief	Coalesce to latest relevant local-day run; regenerate fresh deterministic facts	Label late; no old morning push after the configured usefulness window
Evening check-out	Offer one most recent reconciliation	Do not auto-complete or replay every missed evening
Goal/overdue checks	Evaluate current conditions once on recovery	Suppress conditions no longer true; preserve meaningful unresolved findings
Meeting dossier	Run only before meeting within configured lead/usefulness window	After start, skip prep delivery; use follow-through when notes are available

Profile type	Missed-run rule	Obsolescence rule
Meeting follow-through	Process each unprocessed eligible note revision within configured retention	Revisions update candidates/deduplicate rather than duplicate accepted work
Calendar rescue	Check current availability and pending proposals	Expired slots/conflicts are not actionable; replace affected proposals
Weekly modes	Coalesce per review period and mode	Do not send a stack of historical summaries; history remains inspectable

Exact brief times, lead windows, catch-up maximums, DST policy and retention are configuration gates before launch. A reasonable proposed DST default is once per intended local run, shifting nonexistent local times forward and not duplicating ambiguous-time runs; validate and approve the policy before implementation.

Notification contract

In-app feed is the durable record. Web push is the first out-of-app channel when permitted/available. Email is an opt-in fallback for eligible items when push is unavailable or delivery fails; lack of read acknowledgment alone does not trigger duplicate email. A push-service acceptance is not proof of human receipt. Delivery channels share a stable notification identity.

Tiers: quiet (feed only), scheduled (digest window), nudge (actionable time-sensitive finding), urgent (user-enabled urgent condition). Quiet hours and daily interruption limits apply; urgent does not bypass them unless the user explicitly configures that exception. Coalesce related findings and preserve pending decisions after notification suppression. Email defaults to a minimal summary/link rather than sensitive source content.

Manual app access, feed reading, approval and emergency device unblock never depend on AI budget. Notification transport failure is separate from run success and remains visible.

Shared context & AI budgets

One ingestion layer

Source adapters sync Calendar changes, Gmail history and supplied meeting notes into an owner-scoped store of people, threads, meetings, commitments, goals, summaries and evidence references. Routines query this store, not providers directly. Calendar cursors and Gmail history IDs are opaque provider state with explicit recovery on expiry. A connector may require fetching changed records; incremental does not mean zero network traffic.

First connections use declared bounded backfill (proposed mail starting point: ten days, configurable before consent). Show coverage dates and unknown older promises; do not present bounded history as a complete account. Resync after cursor expiry has its own bounded recovery policy and must preserve accepted local commitments and provenance.

Incremental summaries and freshness

Hash content with owner/source identity and normalization; re-extract only changed material. Version rolling thread/person/goal summaries with source IDs, source versions, model/prompt version and as-of time. Source edits, deletion, permission changes and retention invalidate relevant summaries/context packs. A language-model summary is derived evidence, not ground truth; retrieve specific source evidence for uncertain or consequential claims.

Time passing is an input. Due dates, waiting ages, current local day, meeting lead time, resolved status and capacity are re-evaluated deterministically even when source content hashes are unchanged. Never reuse yesterday's brief as today's factual analysis. If nothing relevant changed, create a short fresh digest with refreshed time-sensitive fields; reuse only still-valid prose fragments, clearly dated.

Deterministic first, language second

Work	Implementation	Guardrail
Overdue, days waiting, double bookings, ranking, capacity, velocity math	Plain code over typed facts and declared preferences	No LLM-computed authority or ungrounded social importance
Commitment extraction and text classification	Configured economical model tier	Schema validation, evidence spans, uncertainty, candidate-only output
Drafts, briefing prose and coaching	Configured stronger model tier on allowed template paths	Grounded context pack; cannot invent facts or choose tools outside schema

No model brand, model identifier or guaranteed price is assumed. Model selection and quality evaluation are deployment configuration. Classification/extraction may be probabilistic even though subsequent checks are deterministic; evaluate errors and keep corrections/audit history.

Context packs and bounded orchestration

Each template profile defines a context recipe, permitted fields, evidence IDs, freshness/coverage requirements, model tier, maximum tokens and maximum steps/tool calls. The orchestrator builds one compact structured pack per logical step/snapshot; sub-steps share its facts and pass typed results, not repeated transcript dumps. Do not reuse packs across users or beyond consent boundaries. Multi-agent orchestration is optional implementation machinery, not a requirement to run several agents for every routine.

Separate trusted template instructions from untrusted source text and free-text preferences. Output style preferences cannot change criteria or authority. Stable profile/preferences prefixes may use provider caching where supported; measure actual cache use rather than assuming savings. Raw credentials, unnecessary source content and secret-bearing fields are excluded.

Budgets and run ledger

Per-user monthly allocation splits scheduled loops, extraction, drafting and chat, with a reserved daily-brief share inaccessible to chat. Add per-loop/run token/step/tool caps, concurrency limits, and a global daily AI kill switch. Conservative numeric limits must be configured before the first pilot; tune after one week of measured use rather than leaving execution uncapped while collecting data.

Reserve maximum allowed usage transactionally before work; reconcile actual known usage after completion. Unknown/cancelled provider usage retains a bounded reservation until reconciliation/expiry policy resolves it. Track model ID, tier, prompt version, input/output/cached tokens, measured/estimated/unknown usage, currency/pricing version, cost estimate, purpose, loop/run IDs and time. Distinguish application allowance from provider invoice cost. Price snapshots and unavailable pricing are explicit.

Usage meter shows purpose breakdown and one proposed 80% threshold alert per budget period (exact threshold configurable). Reserved briefing capacity is finite: if exhausted, deliver a current deterministic digest rather than claiming AI briefing remains guaranteed.

Degradation ladder

> 1. Full permitted output within reservation.

> 2. Use still-valid cached summaries; no new AI analysis; refresh deterministic facts and show coverage/as-of limitations.

> 3. Fresh deterministic digest, without generated prose, retaining pending decisions and source status.

> 4. Pause lower-priority AI steps/loops and disclose what is paused. Continue essential deterministic checks where possible.

No budget mode bypasses approvals, stale-source protections or device-local release. New items that lack extraction remain visibly unprocessed; existing accepted commitments are still tracked. Failures never get silently summarized as no urgent work.

Privacy lifecycle

Before first release define source-specific retention, user deletion/export, disconnect behavior, encryption, model-sharing consent and audit minimization. Removal/permission revocation propagates to raw content, summaries, pack caches and embeddings if used. Retain only minimal non-content audit data permitted by the

agreed retention policy. Existing coach restrictions on raw calendar content are not silently lifted: new context access requires explicitly designed and consented sharing.

Verified technical references

- Google Calendar incremental synchronization: complete pagination before committing the next sync token; invalidated tokens require provider-cache rebuild.
- Gmail synchronization: expired/out-of-range history requires resync; preserve bounded coverage and accepted local work.
- Gmail push: Pub/Sub notifications and watch renewal are connector responsibilities, not execution of user loops.

Implementation & release gates

Status: proposed plan, not authorization to implement or deploy. The first public web release is the complete five-engine executive assistant. Internal build increments are not standalone planner releases.

Existing foundation and scope distinction

Repository documentation describes React/Vite/TypeScript, FullCalendar, FastAPI/Pydantic, PostgreSQL, Celery/Redis, goals/tasks, recurrence, Google integration, proposal review, focus logging, coaching and a Chrome extension. Reuse after targeted assessment, not a rewrite. Existing documented tests do not establish new feature/provider readiness.

A legacy Linux helper in this repository suspends configured executables; it is outside the newly agreed product scope. The new roadmap excludes native-app blocking. Do not claim a hosts-file helper already exists or silently ship legacy app suspension as the new feature. This documentation pass does not remove or alter existing code.

Logical modules

Five domain engines share ingestion/context, proposal/approval, background scheduling and delivery infrastructure. Engines are ownership boundaries in the design, not five required microservices. Keep one maintainable modular application unless scaling evidence justifies separation.

Shared ingestion updates owner-scoped source records and rolling summaries. Commitment reconciliation produces candidates and evidence links. Guardian validates scheduling deterministically. Briefing composes engine outputs. Goal engine tracks outcomes/plans. Insights aggregates coverage-aware statistics. AI sees minimal authorized context packs and cannot execute effects. Template registry constrains trigger/check/output combinations and authority.

Proposed durable records

Record	Required contract
Source record/cursor	Owner, provider identity, source version/hash, sync cursor, coverage/freshness, permissions, tombstone
Summary/context	Subject/source versions, evidence IDs, prompt/model version, as-of, expiry and permission boundary
Commitment/task	Separate obligation and fulfillment task, accepted/candidate state, owner/beneficiary, due/review time, evidence and resolution
Plan/block	Goal/season link, version, accepted status, fixed-time flag, capacity/availability snapshot
Template/loop	Registered schema/version, profile, validated parameters, authority ceiling, context recipe, model tiers/caps, catch-up/stop rules

Record	Required contract
Loop run	Logical trigger key, source/time snapshot, state, budget reservation, output/evidence, sanitized errors
Proposal/approval	Typed effects, before/after, dependencies, versions/destinations, exact authorization, dismissal signature, partial results
Provider action	Stable operation ID, remote version/reference, pending/success/failed/unknown, reconciliation result
Delivery	Stable identity, tier, channels/attempts, consent/quiet-hour/cap decision, feed record
Usage ledger	Purpose/run/model, tokens/cache/estimated cost, reservation/actual/unknown flags, pricing version

Use migrations and owner-scoped indexes/constraints. Decide normalization for high-volume source/run data rather than extending planning JSON indiscriminately. Preserve accepted data and evidence during recovery/migration.

Internal build order

Stage	Work	Exit evidence
0	Finalize templates, UX, retention/consent, limits and scope	Walkthrough of full assistant journey; built-in check schemas validated and unresolved operational defaults approved
1	Calendar, goals and tasks; calendar-first shell; guardian/proposal foundation	Google read/sync/publication acceptance; fixed blocks, stale source, retry, conflicting/partial write checks
2	Shared context, registry, server loops, budgets and delivery; focus web flow	Closed-app brief, catch-up, dedup, costs, quiet hours, Chrome website expiry, timer/audio recovery
3	Meetings from supplied notes; commitment ledger and dossiers	Extraction/merge/resolution quality; provenance; task/draft approval; imported history coverage
4	Gmail; email commitments, reply/waiting/VIP, noise insights	Provider consent/verification, bounded ingestion, expired history recovery, approved draft/send, privacy lifecycle
5	Full engine integration and pilot	All F01-F21/N01-N08 evidence; real-use pilot and tuning before release decision
Later	Desktop timers for macOS/Windows/Linux; one-OS system website helper; other providers/capture/handoff	Separate platform gates; native-app/cross-device blocking remains excluded

Stage 1/2 prototypes may be reviewed locally, but must not be presented as the requested full release. Meeting notes import satisfies the first release; bot-free recording is not a hidden dependency. Google Calendar is foundational and Gmail is the proposed initial mail provider; additional providers require later decisions.

Consistency and execution

Calendar source events are observations, not proposals: display already-made external changes promptly, then propose Kairos corrections. Protect fixed blocks. Recheck owner, consent, proposal version, source freshness, current availability and dependencies before apply. Persist local intent and use a durable outbox or equivalent for external effects; no distributed transaction is implied. Stable keys and reconciliation prevent duplicate/unknown writes from being blindly retried.

Use server scheduler/worker jobs with leases and dedup, bounded retries and pause/version checks before delivery. Device owns session expiry and emergency unblock. Browser shutdown does not stop server loops and server failure does not prolong website rules. AI-off mode and budget kill switch preserve manual functionality and current deterministic findings.

Acceptance and pilot

Trace every requirement to implementation issue and evidence. Test permission/ownership, unsupported template values, authority escalation, source injection, stale approvals, fixed blocks, dismissal suppression, dependent batch changes, cursor expiry, missed runs, DST/time progression, duplicate delivery, cost reservation races and device recovery. Use labeled examples to measure extraction, unresolved-vs-answered distinctions and grounding, not just schema validity.

Before pilot configure retention and numerical budgets, performance dataset/budgets, notification limits and supported Chrome/OS versions. Measure real usage for at least the proposed first week before tuning caps. Production release additionally needs applicable provider verification, operational monitoring, backup/restore, migration rollback, consent/privacy review and a separate deployment authorization. No tests, deployment or live provider changes are claimed by this documentation update.

Verified platform dependencies

Google Calendar sync uses paginated changes and invalidated-token recovery. Calendar push requires a public HTTPS receiver; notifications signal the need to synchronize rather than carrying the complete changed event. See Calendar push. Polling can support local validation; closed-app production automation requires hosted workers/scheduler and monitoring, subject to later deployment authorization.

Gmail sync requires history-expiry recovery. Gmail push requires Pub/Sub and renewable watches; configure renewal and polling/recovery fallback. Gmail scope documentation identifies restricted scopes and applicable verification/security-assessment dependencies for server-handled restricted data; validate audience, scopes and exceptions before setting a public launch date.

Chrome declarative request rules apply within supported browser/profile capability, not native-app control. Web push overview explains subscription/service-worker delivery; server jobs remain authoritative. Document Picture-in-Picture may be an optional browser enhancement but does not replace the later cross-platform desktop companion contract.

Technical references checked 30 October 2026. These constraints inform implementation planning, not a claim that new connectors or desktop packages already exist.

Decisions & references

Version 2.1 / 31 October 2026. Latest explicit user decisions override older assumptions. Product scope is separate from current repository capability.

Confirmed decisions

ID	Decision
D01	Full AI executive-assistant first web release: commitment ledger, calendar guardian, briefing composer, goal engine and insights. Not a planner-only MVP.
D02	Google Calendar reading, synchronization and approved publication are foundation; other calendar providers later.
D03	Fixed declarative templates, configured in one short chat with preview/confirmation. Parameters cannot raise authority; style text changes writing only; unsupported requests are logged and nearest supported option offered.
D04	Six core families: Daily brief, Daily check-out, Goal check-in, Overdue and slipping work, Focus rescue, Weekly review. Friday wrap and week-ahead are shared-data modes.
D05	Daily coaching preserves accepted blocks; propose only justified changes. Keep this time fixed protects selected blocks.
D06	Commitment resolution uses evidence; reminders, replies, sessions and task counts alone are not proof of fulfilled outcomes.
D07	In-app timer first; desktop timer companions follow for macOS, Windows and Linux.
D08	Chrome website blocking in v1. Later system-wide website blocking on one selected OS. No native-app or cross-device blocking promises.
D09	Related changes bundle with dependencies; obsolete proposals are replaced; equivalent dismissed proposals remain suppressed until meaningful facts change.
D10	Server-side loops, catch-up policies and quiet/scheduled/nudge/urgent notification tiers; quiet hours and daily limit. In-app feed, then web push, then opt-in email fallback.
D11	Shared incremental source/context store; rolling summaries, bounded backfill, deterministic checks, bounded source-linked packs, economical extraction tier and stronger writing tier.

ID	Decision
D12	Purpose-specific budgets with brief reservation, per-run/loop limits, kill switch, usage/cost ledger and graceful degradation; values tuned from usage.
D13	Web first, desktop later. Internal build order calendar/goals/tasks, meetings, email; full assistant release gate includes all five engines.
D14	Today, Plan, Commitments, Goals, Review, Assistant; Today/Plan retain the calendar-first workspace and collapsible coach panel.
D15	Single purple accent with white/light gray, charcoal text and subtle borders. Earlier pastel palette is superseded everywhere.
D16	Detailed documentation and updated PDF now; no app-code implementation or deployment in this task.
D17	Confirmed 31 October 2026: specialized assistant workflows are built-in checks/output profiles inside the six templates, not additional selectable templates.

Assessment and proposed resolutions

The five domain engines supply capabilities; the six templates expose those capabilities through registered built-in checks/output profiles. The user confirmed this architecture on 31 October 2026. The specialized routines remain fully in scope without adding top-level workflow templates. Automation defines the mapping for schema and workflow validation.

Gmail is the proposed initial mail provider, consistent with Calendar's Google foundation and the requested build order. Meetings initially use supplied notes/transcripts rather than native capture. Native capture, other calendar/mail providers, relationship-cooling/Going cold and AI handoff are later additions unless separately prioritized.

The legacy Linux implementation actually suspends apps; it is not a hosts helper. New product scope excludes that capability. Existing code is neither silently removed nor reclassified as shipped website blocking.

Today's coaching is interactive use of Daily brief, not a seventh automation template. A generated report or preview can appear automatically; accepting an inferred task or applying an external draft/calendar/send effect needs scoped authority. Calendar external edits are observed facts, not retroactively gated by approval.

Open decisions and release gates

Item	Proposed handling / decision needed	Gate
First desktop blocking OS	Choose macOS, Windows or Linux; timer support for all three does not imply website-helper parity	Later desktop helper
GPS framework	Confirm exact method/wording before claiming faithful named-method support	Goal coaching content

Item	Proposed handling / decision needed	Gate
Audio	Select licensed source/tracks and offline/playback policy	Focus acceptance
Session recovery	Define sleep/restart/abandonment/overrun rules and fixed expiry behavior	Timer/device implementation
Loop defaults	Set brief times, lead/catch-up windows, quiet hours, daily caps and DST policy	Activation defaults
Source/privacy	Approve source-specific retention, model sharing, disconnect/delete/export and bounded backfill	First ingestion pilot
Budgets/quality	Conservative numeric caps before pilot; tune from real usage; define labeled evaluation and UX performance targets	Pilot
Gmail launch readiness	Provider verification/security obligations and live acceptance depend on selected permissions and audience	Public release

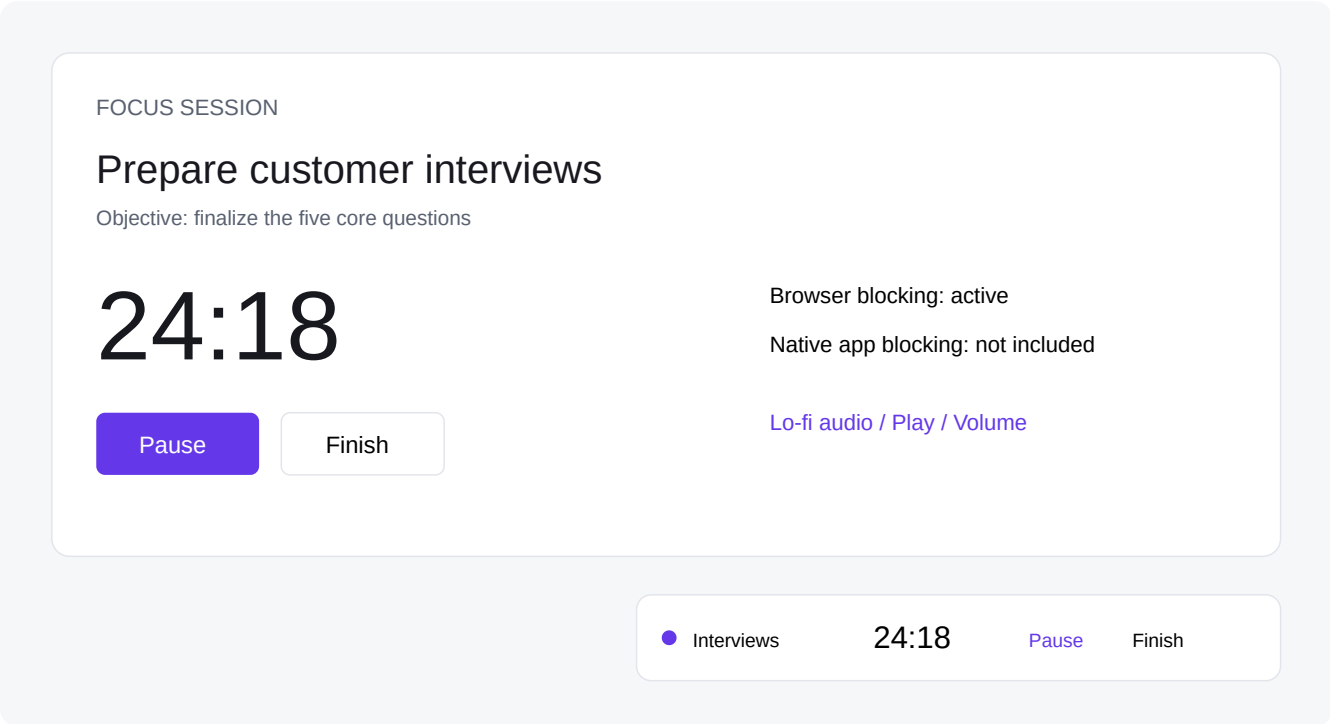
References and provenance

User messages establish decisions. Screenshots/reference notes inform hierarchy and interactions, not independent instructions to execute. Shared ChatGPT link was unavailable and is not reviewed evidence. Sources referenced for product inspiration: Superfocus, Kai, Kai workflows, Kai integrations. Organization references: Notion product specs and Jama PRD guide.

Technical references are collected in Context and Implementation as verified during this revision. Marketing illustrations do not prove connector or platform capabilities. PDF chapters are generated from the Markdown without hidden product wording substitutions.

FOCUS CONCEPT / S07-S08

One task. Clear controls.



Conceptual workstation and floating in-app timer. Report actual device coverage; sound is optional. Desktop timer companions for macOS, Windows and Linux follow the web release. System website blocking starts later on one selected OS.

Completion is a separate decision

Finish the session, record the accomplishment, explicitly update task status and optionally check in on an outcome. Unfinished work can produce a new proposal; it never silently reschedules itself.